



**Barcelona
Supercomputing
Center**

Centro Nacional de Supercomputación



Practical Persistent Memory Programming

Adrian Jackson

EPCC

a.jackson@epcc.ed.ac.uk

@adrianjhpc 

Javier Conejero

BSC

francisco.conejero@bsc.es

Aims

- Understand persistent memory hardware and software
- Understand the different aspects that can impact performance, and the shared/private nature of the resources
- Learn how to program persistent memory
- Get hands on with persistent memory hardware



Aims cont

- Understand data movement and think about application data requirements
- Understanding I/O and data movement, particularly persistence requirements, is hard
- Thinking about different ways you can do I/O or storing data

Format



- Lectures and practicals
- Slides and exercise material available online:
 - <https://github.com/NGIOproject/PMTutorial>
 - Exercises will be done on remote machine (NEXTGenIO prototype)
 - We will give you accounts on these

Timetable

- 08.30 Introduction
- 08.45 Practical: IOR and Streams
- 09.00 Hardware, I/O and storage
- 09.30 Practical: Using different mount points
- 09.45 Persistent thinking
- 10.00 Break
- 10.30 Low-level persistent memory programming
- 10.45 Practical: Persistent memory programming
- 11.00 Higher-level persistent memory programming
- 11.15 Persistent Memory Programming with PyCOMPSs
- 12.00 Finish

Programming persistent memory



SC19 Tutorial: Practical Persistent Memory Programming

Programming persistent memory



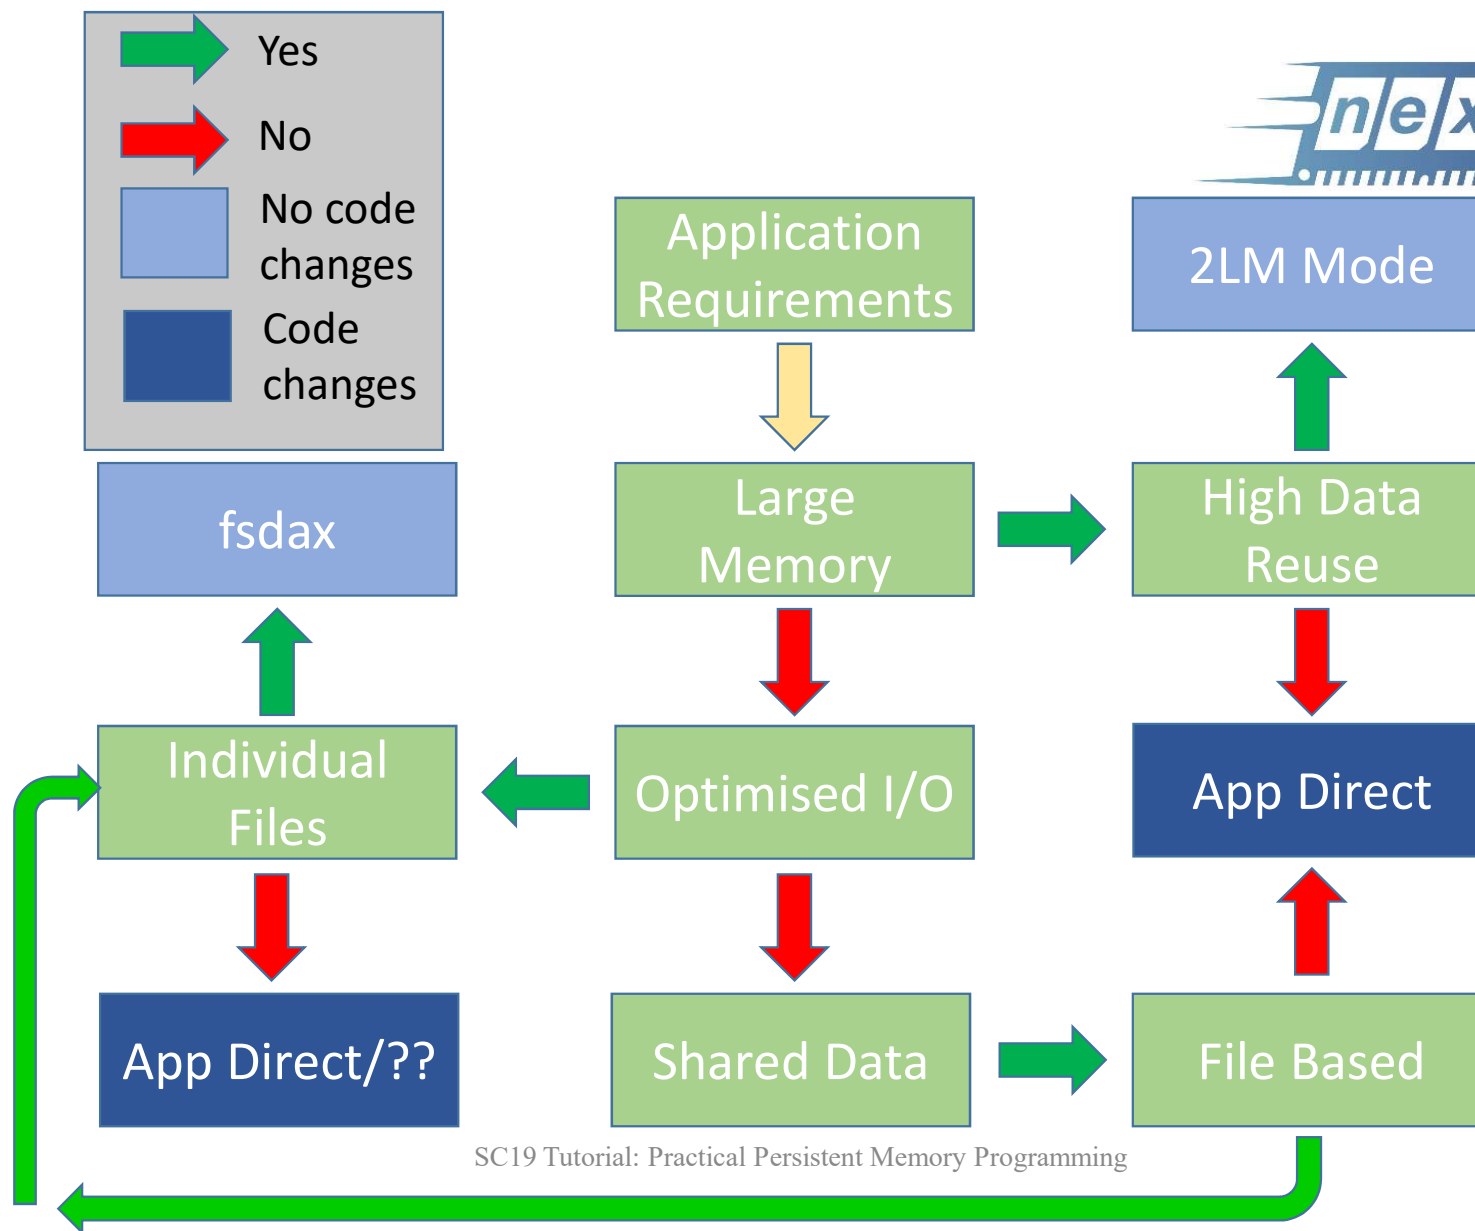
```
double *a, *b, *c;
pmemaddr = pmem_map_file(path, array_length,
                        PMEM_FILE_CREATE|PMEM_FILE_EXCL,
                        0666, &mapped_len, &is_pmem)

a = pmemaddr;
b = pmemaddr + (*array_size+OFFSET)*BytesPerWord;
c = pmemaddr + (*array_size+OFFSET)*BytesPerWord*2;

#pragma omp parallel for
for (j=0; j<*array_size; j++){
    a[j] = b[j]+scalar*c[j];
}
pmem_persist(a, *array_size*BytesPerWord);
```

Programming persistent memory





Programming persistent memory



- Design and performance considerations are the challenge
 - Programming the memory is easy
- Design for functionality
 - What is persistent, when is it persistent, what failures can you tolerate, etc..
- Design for performance
 - Memory size, I/O, data access costs, etc...
- Design for hardware configurations
 - NUMA, filesystems, storage, etc...

Summary



- Please don't hesitate to ask questions!
- We are doing practicals
 - But if you're not confident in programming we have other options
- We are aiming at different experience levels so if it's too easy/you know it already/it's too difficult let us know

Practical



- Take IOR and STREAMS source code

- Run on the prototype

- Prototype is available at:

```
ssh hydra-vpn.epcc.ed.ac.uk
```

- Then

```
ssh nextgenio-login2
```

- Using your guest account

- Practical source code is available at:

```
/home/nx01/shared/pmtutorial/exercises
```

- Using slurm as the batch system

- pm partition for this course

```
sbatch -p pm scriptname.sh
```



ssh [nggquest01@hydra-vpn.epcc.ed.ac.uk](ssh:nggquest01@hydra-vpn.epcc.ed.ac.uk)
ssh nextgenio-login2