



Persistent Memory Programming

Adrian Jackson

EPCC

a.jackson@epcc.ed.ac.uk

@adrianjhpc

Javier Conejero

BSC

Christian Herold

TU Dresden

Aims



- Understand persistent memory hardware and software
- Understand the different aspects that can impact performance, and the share nature of the resources
- Learn how to program persistent memory
- Understand profiling and analysing persistent memory usage
- Get hands on with persistent memory hardware

Real Aims



- Understand data movement and application data requirements
- Understanding I/O and data movement is hard
- Thinking about different ways you can do I/O or storing data
- Understand some tools that can help you

Format



- Lectures and practicals
- Slides and exercise material available online:
 - <https://github.com/adrianjhpc/iotutorial>
- Exercises will be done on remote machine
 - NEXTGenIO prototype
 - We will give you accounts on these
- Some programming for those who want to, for others can just run programs
- Some setting up of profiling tools

Timetable




- 09.00 Persistent memory hardware
- 09.30 Persistent memory performance practical
- 10.00 Persistent memory software
- 10.45 Break
- 11.30 Persistent memory programming practical
- 11.30 Alternative methods of exploiting persistent memory
- 12.00 Profiling and debugging persistent memory applications
- 12.30 Summary and Finish

New Memory Hierarchies



- High bandwidth, on processor memory
 - Large, high bandwidth cache
 - Latency cost for individual access may be an issue
- Main memory
 - DRAM
 - Costly in terms of energy, potential for lower latencies than high bandwidth memory
- NVRAM main memory
 - High capacity, ultra fast storage
 - Low energy (when at rest) but still slower than DRAM



Non-volatile memory







- Non-volatile RAM
 - 3D XPoint technology
 - STT-RAM
- Much larger capacity than DRAM
 - Hosted in the DRAM slots, controlled by a standard memory controller
- Slower than DRAM by a small factor, but significantly faster than SSDs
- STT-RAM
 - Read fast and low energy
 - Write slow and high energy
 - Trade off between durability and performance
 - Can sacrifice data persistence for faster writes

SRAM vs NVRAM



- SRAM used for cache
- High performance but costly
 - Die area
 - Energy leakage
- DRAM lower cost but lower performance
 - Higher power/refresh requirement
- NVRAM technologies offer
 - Much smaller implementation area
 - No refresh/ no/low energy leakage
 - Independent read/write cycles



NEXTGenIO - key facts



- Research & Innovation Action
- 48 month duration
- €8.1 million
- Approx. 50% committed to hardware development
- Prototype system part of the project



Project objectives






- Hardware platform prototype
 - Demonstrating the prototype's broad applicability for both HPC and data centric applications
- Exascale I/O investigation
 - Understanding how best to exploit NVRAM
- Systemware development:
 - Producing the necessary software to enable Exascale application execution on the hardware platform
- Application co-design
 - Understanding individual application I/O profiles and typical I/O workloads on shared systems running multiple different applications

Hardware



- Motherboard developed at Fujitsu factory in Augsburg





ISC19: Persistent Memory Programming

Systemware





- Work on adapting job scheduler (SLURM)
- Development of a data scheduler
- Object stores as alternatives to file systems
 - DAOS (Distributed Application Object Storage)
 - dataClay
- Multi-node NVRAM file system
 - echoFS


Burst Buffer



- Non-volatile already becoming part of HPC hardware stack
- SSDs offer high I/O performance but at a cost
 - How to utilise in large scale systems?
- Burst-buffer hardware accelerating parallel filesystem
 - Cray DataWarp
 - DDN IME (Infinite Memory Engine)

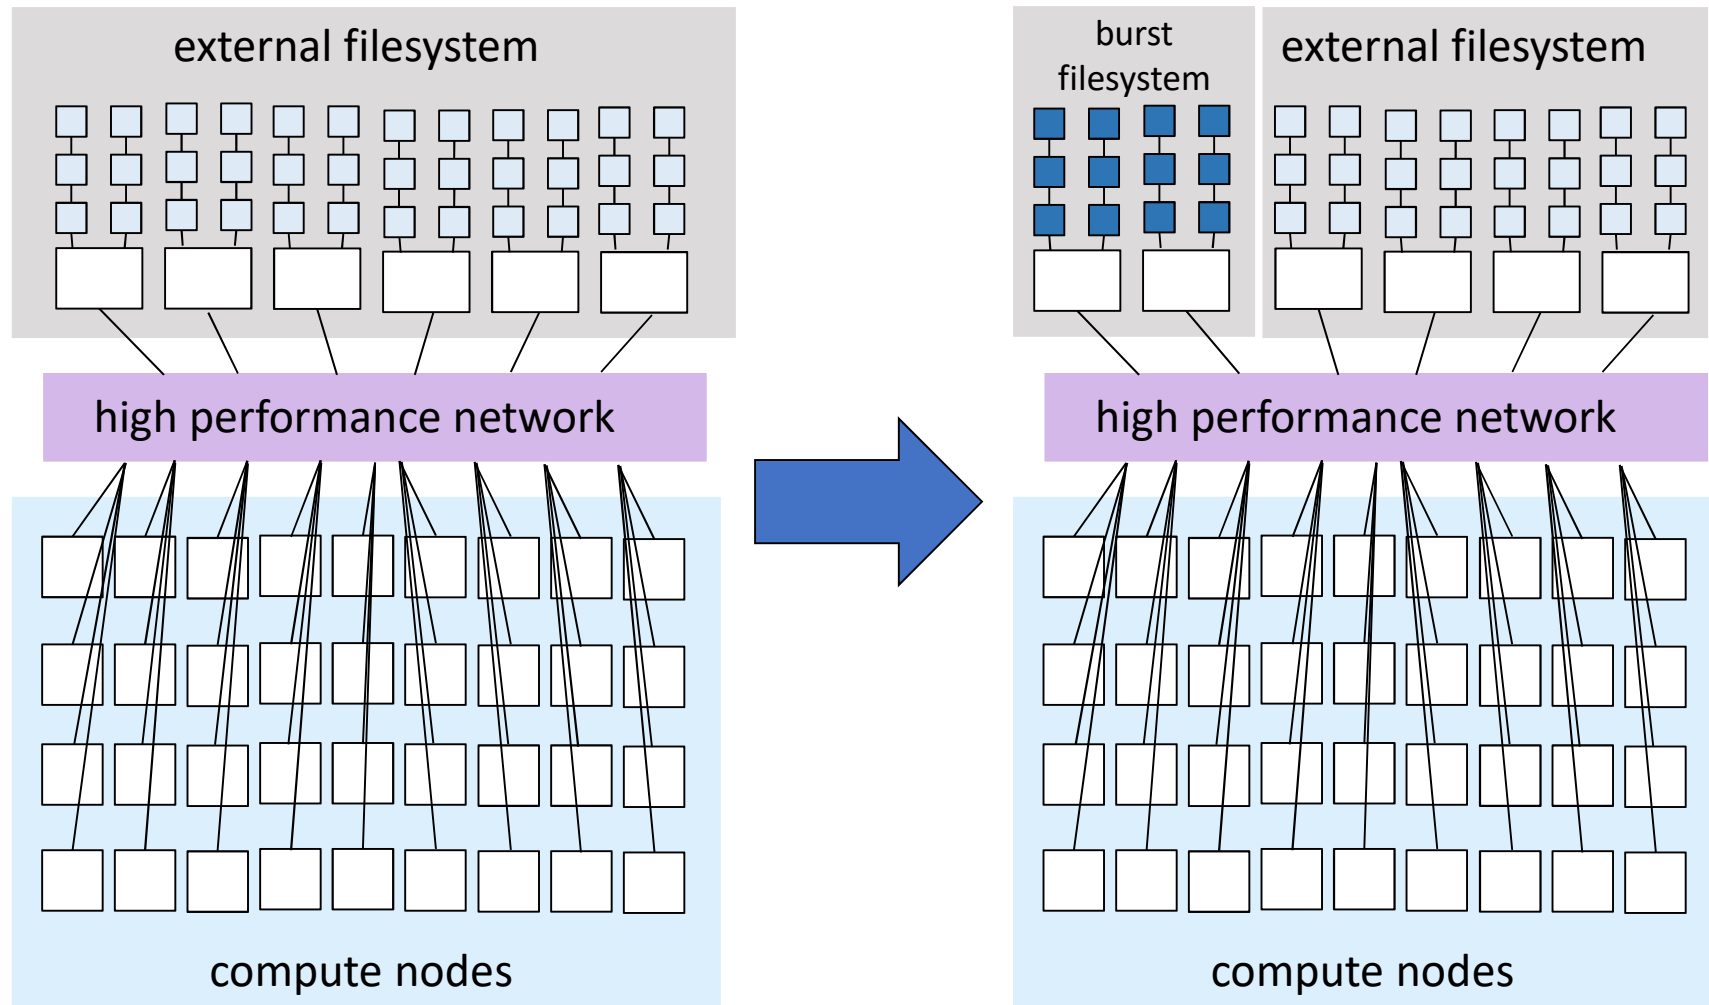


CRAY



Sec19: Persistent Memory Programming


Burst buffer



Burst buffer



- Burst buffer is (generally) separate resource to filesystem
- Managing usage/allocation/data movement is user responsibility
- Storage has compute
- Resource is external
- Scheduling can be separate
 - Usage not required



Summary



- Please don't hesitate to ask questions!
- We are doing practicals
 - But if you're not confident in programming we have other options
- We are aiming at different experience levels so if it's too easy/you know it already let us know